



CSCI 3411 – OS

Lab 4

SYNCHRONIZATION

Overview

- ▶ Review – Function Pointers and Pthreads
- ▶ Mutex
- ▶ Semaphore
- ▶ Monitor

C Function Pointers

- ▶ Declaration: *returnType* (* ptrName) (*paramList*);
 - ▶ Ex. `int (* myFuncPtr)(int, int);`
 - ▶ Declares a function pointer which can only point to a function that takes two int parameters and returns an int.
- ▶ Assignment: `myFuncPtr = &myFunction;`
- ▶ Dereference: `(* myFuncPtr)(5, 11);`

C Function Pointers

- ▶ Why can't you declare a function pointer like these:
 - ▶ `int *myFuncPtr(int, int);`
 - ▶ `int *(myFuncPtr(int, int));`
- ▶ Why can't you deference the function pointer:
 - ▶ `*myFuncPtr(0, 10);`
- ▶ Answer: C's Operator Precedence
 - ▶ Function Calls (i.e. (...)) are evaluated before * operator.
 - ▶ * operator binds to the type by default, not the variable name.

Pthreads

- ▶ Threading mechanism for C (POSIX threads).
- ▶ Usage:
 1. Create variables to hold thread ID numbers.
 2. Create thread attributes (can leave at defaults – priority level, etc.)
 3. Create and start threads:
 - ▶ Supply what function each thread should run and what arguments to pass to these functions.
 4. (Optionally) let the parent wait until all threads are finished.

Pthreads

► Usage:

1. Include `<pthread.h>`
2. Create `pthread_t` variables to hold ID numbers (one for each thread).
3. Create `pthread_attr_t` variable and set it to default values.
 - `pthread_attr_init(pthread_attr_t *)`
4. Create and start threads:
 - `pthread_create`
`(pthread_t *, pthread_attr_t *, void *(* name)(void *), void *)`
 - In order: variable for ID number, attributes, function to execute, arguments for the function.
5. (Optional) let the parent wait until all threads are finished.
 - `pthread_join(pthread_t, NULL);`

Mutex

- ▶ Mutual Exclusion
- ▶ Locking mechanism – `acquire()` and `release()`
 - ▶ Built-in “boolean” variable.
- ▶ “Waiting” is implemented via a busy loop / spinlock.
 - ▶ Effective if the time to context switch > busy waiting.

Mutex

► Usage:

1. Call `acquire()` to obtain the mutex lock.
 - If someone else has it, wait.
2. Execute code in critical section.
3. Call `release()`

```
mutex.acquire();  
/* Critical section code */  
mutex.release();
```


Mutex in C

- ▶ Included in `<pthread.h>`
- ▶ `pthread_mutex_t` - defines a mutex type
- ▶ `pthread_mutex_init (pthread_mutex_t *, pthread_mutex_attr *)`
 - ▶ Call to initialize a mutex. The second param can be NULL => defaults.
- ▶ `pthread_mutex_lock (pthread_mutex_t *)`
 - ▶ Same as `acquire()`
- ▶ `pthread_mutex_unlock (pthread_mutex_t *)`
 - ▶ Same as `release()`
- ▶ `pthread_mutex_destroy (pthread_mutex_t)`

Mutex in C - Usage

```
#include <pthread.h>

pthread_mutex_t mutex;

int main(){
    pthread_mutex_init(&mutex);
    ...
    // Create, run, and join threads
    ...
    pthread_mutex_destroy(&mutex);
}

void *threadFunc(void *params){
    pthread_mutex_lock(&mutex);

    // Critical section

    pthread_mutex_unlock(&mutex);
}
```

Semaphores

- ▶ Can be used to control resource access for N resources.
 - ▶ Example: a system to reserve study rooms in the library.
- ▶ `wait()` and `signal()`.
- ▶ Study room semaphore initialized to 10
 - ▶ Calling `wait` decrements the value by 1, unless there are none left, in which case you wait for one to become available.
 - ▶ Calling `signal` increases the value by 1.
 - ▶ As opposed to having one mutex per room.
- ▶ A semaphore initialized to 1 == mutex.

Semaphores

► Usage:

1. Initialize a semaphore to N (to represent N resources or N allowable accesses).
2. Call `wait()` to decrement a semaphore's value by one.
 - If ≤ 0 , wait until the resource is available.
3. Execute critical section.
4. Call `signal()` to increment a semaphore's value by one.

Semaphore Implementation

- ▶ Wait can be implemented with a waiting loop, like the mutex.
- ▶ Alternative to waiting loop:
 - ▶ Each semaphore keeps track of a process queue.
 - ▶ In `wait()` – instead of causing a process to spinloop, add it to the process queue and suspend its execution.
 - ▶ In `signal()` – wakeup the process at the head of the queue.

Semaphores in C

- ▶ Included in `<semaphore.h>`
- ▶ `sem_t` – defines a semaphore type
- ▶ `sem_init (sem_t *, int, int)`
 - ▶ The first `int` is a flag to set if it should be shared with a forked process.
 - ▶ The second is the semaphore's initial value.
- ▶ `sem_wait (sem_t *)`
- ▶ `sem_post (sem_t *)`
 - ▶ Same as `signal()`
- ▶ `sem_destroy (sem_t *)`
- ▶ `sem_value (sem_t *, int *)`
 - ▶ Retrieves the current value of the semaphore.

Semaphores in C - Usage

```
#include <semaphore.h>

sem_t semaphore;

int main(){
    sem_init(&semaphore, 0, 5);
    ...
    // Create, run, and join threads
    ...
    sem_destroy(&semaphore);
}

void *threadFunc(void *params){
    sem_wait(&semaphore);

    // Critical section

    sem_signal(&semaphore);
}
```

Monitors

- ▶ Encapsulates programmer-defined functions and variables into one synchronized data type (a module, class, object, etc.).
- ▶ Special “condition” variables that can be used with `wait()` or `signal()`
- ▶ Only one process is allowed at a time within a monitor.
- ▶ Multiple processes can be waiting in a queue to enter or waiting on specific condition variables.

Monitors

```
monitor myMonitor {  
    int localVar1;  
    char localArr1[];  
    condition x, y, z;
```

```
    void fun1() {  
        ...  
        x.wait();  
        y.signal();  
    }
```

```
    void fun2 () {  
        y.wait();  
        ...  
    }
```

```
    void fun3 () {  
        x.signal();  
        ...  
    }  
}
```

Usage

Thread 1:

```
myMonitor.fun1(); // Waiting on x  
printf("A\n");
```

Thread 2:

```
myMonitor.fun2(); // Waiting on y  
printf("B\n");
```

Thread 3:

```
myMonitor.fun3(); // Signals X  
printf("C\n");
```

Monitors

- ▶ Differences from other synchronization constructs:
 - ▶ A monitor includes synchronized functions and variables.
 - ▶ Only one process allowed inside at a time (unlike a semaphore).
 - ▶ If nobody is waiting on a condition, calling its `signal()` doesn't do anything.
- ▶ Less room for erroneously calling `wait` / `signal`.
 - ▶ With semaphores:
 - ▶ Calling `signal(...)` before `wait(...)`
 - ▶ Calling `wait(...)` and then `wait(...)` again instead of `signal(...)`
 - ▶ Omitting `wait(...)` or `signal(...)`

Monitors

- ▶ Process P calls `signal()` is called on a condition variable which Process Q is waiting on. When do you let Q inside the monitor?
 - ▶ Signal and wait: Q is let inside immediately and P waits until Q leaves (or waits).
 - ▶ Signal and continue: Q waits until P leaves (or waits).

Monitor Implementation

- ▶ Monitors are really only a concept (an abstract data type). At the end of the day, you need some way to provide the synchronization features....
 - ▶ Implementation via semaphores (pg. 229):
 1. Semaphore to guard entrance to each function.
 2. Semaphore(s) to guard access to each condition variable.
 3. Semaphore used to implement “signal and wait”.

Questions?